

University of Waterloo E-Thesis Template for L^AT_EX

by

Pat Neugraad

A thesis
presented to the University of Waterloo
in fulfillment of the
thesis requirement for the degree of
Doctor of Philosophy
in
Philosophy of Zoology

Waterloo, Ontario, Canada, 2023

© Pat Neugraad 2023

Examining Committee Membership

The following served on the Examining Committee for this thesis. The decision of the Examining Committee is by majority vote.

External Examiner: Bruce Bruce
Professor, Dept. of Philosophy of Zoology, University of Wallamaloo

Supervisor(s): Ann Elk
Professor, Dept. of Zoology, University of Waterloo
Andrea Anaconda
Professor Emeritus, Dept. of Zoology, University of Waterloo

Internal Member: Pamela Python
Professor, Dept. of Zoology, University of Waterloo

Internal-External Member: Meta Meta
Professor, Dept. of Philosophy, University of Waterloo

Other Member(s): Leeping Fang
Professor, Dept. of Fine Art, University of Waterloo

Author's Declaration

I hereby declare that I am the sole author of this thesis. This is a true copy of the thesis, including any required final revisions, as accepted by my examiners.

I understand that my thesis may be made electronically available to the public.

Abstract

This is the abstract.

Acknowledgements

I would like to thank all the little people who made this thesis possible.

Dedication

This is dedicated to the one I love.

Table of Contents

Examining Committee	ii
Author’s Declaration	iii
Abstract	iv
Acknowledgements	v
Dedication	vi
List of Figures	viii
List of Tables	ix
1 Implementation and benchmarking	1
1.1 WolfSSL	1
1.1.1 Integrating post-quantum algorithms	1
1.1.2 Implementing KEMTLS	2
1.1.3 Miscellaneous comments	4
References	5
APPENDICES	6

List of Figures

List of Tables

Chapter 1

Implementation and benchmarking

To evaluate post-quantum TLS 1.3 and KEMTLS on constrained devices, we implemented post-quantum TLS 1.3 and KEMTLS on top of WolfSSL, a TLS library written in the C programming language. We then implemented a minimal TLS client on the Raspberry Pi Pico 2 W, a microcontroller with two ARM Cortex-M33 cores ¹ and 512kB of SRAM, and measures the time it takes for the Pico client to complete a TLS 1.3 or KEMTLS handshake with a server. This chapter describes some of the implementation details, the benchmarking methodology, and the performance measurements.

1.1 WolfSSL

WolfSSL is a modern open-source TLS library written in C. In addition to a complete TLS stack, WolfSSL also includes its own cryptography library called WolfCrypt. Both WolfSSL and WolfCrypt are optimized for code size, speed, and memory footprint, and its portability and ease of configuration greatly simplifies managing multiple build targets using a single code base.

1.1.1 Integrating post-quantum algorithms

As of June 2025, WolfCrypt contains an in-house implementation of ML-KEM and ML-DSA. Both implementation are skillfully optimized, achieving at least 2x speedup on the

¹The Pico 2 W also has 2 RISC-V cores, though we did not use them in this project.

Pico compared to the reference implementation ². Unfortunately, this makes the comparison across different schemes unfair. Instead, I chose to integrate with PQClean’s clean implementations.

While PQClean is not specifically optimized for embedded system builds, all of its clean implementations are trivially portable to ARM. One non-trivial challenge is adapting the `randombytes` API in PQClean to an embedded system build with no operating system. Fortunately, WolfCrypt’s `WC_RNG` API provides a common abstraction that works on both a desktop build (where random bits can be sourced from `/dev/urandom`) and Pico (where random bits are collected from various peripherals by the SDK). In the end, we expanded the `randombytes` API so it can be told to source random bits from user specified instance of `WC_RNG`.

Expanding the selection of KEMs for the initial key exchange (i.e. `ClientHello` and `ServerHello`) is straightforward, thanks to the fact that WolfSSL already supports ML-KEM for key exchange. The `NamedGroup` enum is trivially captured using a single 16-bit integer, and the logic for branching into the correct KEM allows for a simple `switch-case` block.

Expanding the selection of post-quantum signatures is trivial thanks to previous efforts to integrate `liboqs` into WolfSSL. We only need to replace all uses of `liboqs` with their equivalents in PQClean.

1.1.2 Implementing KEMTLS

Generating certificate chain and private keys

WolfCrypt’s `asn.h` API provides a nearly complete collection of tools needed to generate certificate chains, encode certificates and private keys according to DER, then further encode them to PEM format. At the time of writing this thesis, WolfCrypt does not support signing a certificate signing request (CSR), but for benchmarking purposes I control the entire chain, and WolfCrypt does support directly signing the body of a certificate.

Modifying WolfCrypt’s `asn.h` module to support KEM public key in a certificate is relatively straightforward. The only non-trivial obstacle comes from how WolfSSL handles OIDs. Object Identifier (OID) is a variable-length sequence of integers used to identify individual cryptographic primitives. For example, the OID for ML-KEM-512 is

²The Keccak implementation in WolfCrypt is only roughly 10% faster than PQClean’s implementation, so the optimization must have come elsewhere

2.16.840.1.101.3.4.4.1. OID is included in an certificate to identify the public key and the signature; it is also included in the DER encoding of private keys. Having variable length makes OID tedious to work with when programming in C: unlike `NamedGroup`, which has fixed length that can be captured in a 16-bit integer, OID cannot be easily abstracted using a fixed-sized enum type. WolfSSL works around this limitation by using an “OID sum” algorithm, which computes an “hash” of an OID that fit into a 32-bit integer. Compressing variable-length integer sequence into a 32-bit integer carries with it the risk of collision, and in fact the first version of the OID summing algorithm indeed ran into a collision between SPHINCS-192-fast and SPHINCS-128-fast. A newer OID summing algorithm provided stronger collision resistance and resolved this issue. All OID sums are stored in a header file `oid_sum.h`, which is generated by a Perl script.

Implementing unilaterally authenticated KEMTLS handshake

KEMTLS handshake workflow is identical to TLS 1.3’s handshake flow from the beginning until client starts processing server’s `Certificate` message. Even after KEMTLS and TLS 1.3’s handshake flow diverges, they still share the format of the `Finished` message (which contains exactly one HMAC tag). Finally, once the handshake is complete, TLS 1.3 and KEMTLS exchange application data in identical fashion. The similarity between TLS 1.3 and KEMTLS handshake workflow allows us to reuse a significant part of WolfSSL’s TLS 1.3 implementation, diverging at only a handful of places that are easy to reason about. While working with a TLS library written in C is intimidating at first, this implementation strategy proved successful, and I was able to finish implementing KEMTLS in less than a month using only around 4600 lines of code change.

The `WOLFSSL` struct is used on both client-side and server-side and encodes the pair of client and server state as a global TLS state. We begin modifying the TLS state machine by adding two flags `haveMlKemAuth` and `haveHqcAuth` to the main `WOLFSSL` struct. In a unilaterally authenticated KEMTLS handshake, the two flags are set on the server side when the server loads a KEM private key. Detecting a KEM private key is cleanly accomplished because at certificate generation, private keys are encoded according to DER, and the OID of the KEM scheme is included. If the OID belongs to one of ML-KEM’s variants, then `haveMlKemAuth` is set, and if the OID belongs to one of HQC’s variants, then `haveHqcAuth` is set. On the client side, these two flags are set when the client finds a ML-KEM or HQC public key in the certificate chain sent by the server. The combination of these two flags is sufficient for deciding when the two peers are performing a KEMTLS or TLS 1.3 handshake, and all divergence between KEMTLS and TLS 1.3 handshake flow will be controlled by these two flags.

On the client side, KEMTLS and TLS 1.3 handshake flows first diverge after the client finishes processing server's **Certificate** message. In signature-based TLS 1.3, client's immediate next step is to receive and process server's **CertificateVerify** containing a signature over the handshake transcript. In KEMTLS, client will not receive additional message. Instead, it uses the KEM public key to encapsulate the **authentication secret**, then sends the ciphertext to the server in a **KemCiphertext** message. We followed the original KEMTLS implementation's **KemCiphertext** format as a handshake message whose payload contains the raw ciphertext and no additional metadata. Within the context of this project, each server instance will only load one private key, so **KemCiphertext** not carrying metadata on the ciphertext will not cause confusion. However, WolfSSL supports loading multiple private keys for authentication, in which case it might be necessary for **KemCiphertext** to carry metadata such as an OID.

Client's **KemCiphertext** is encrypted under handshake traffic keys derived from the unauthenticated handshake secret (HS). After sending **KemCiphertext**, client must update the key schedule by mixing in the authentication secret and deriving the authenticated handshake secret (AHS). From AHS, the client will derive new handshake traffic key for encrypting additional handshake messages, as well as application traffic key. This key schedule update is necessary for subsequent messages to provide implicit authentication: no adversary, even if it compromises the handshake secret, can decrypt subsequent handshake message or application data without the long-term secret key.

Client's **Finished** follows the same format as TLS 1.3's **Finished**: a handshake message whose payload contains a raw HMAC tag computed under the **finished_key** against the handshake transcript. The MAC key is derived from the **MasterSecret**, which is itself derived from AHS. After sending **Finished**, the client can start sending application data without receiving server's **Finished** first. Because application data will be encrypted under application traffic key derived from authenticated handshake secret, any server successfully decrypting them is implicitly authenticated. Finally, client explicitly authenticates the server after receiving server's **Finished**.

1.1.3 Miscellaneous comments

References

- [1] Michel Goossens, Frank Mittelbach, and Alexander Samarin. *The L^AT_EX Companion*. Addison-Wesley, Reading, Massachusetts, 1994.
- [2] Donald Knuth. *The T_EXbook*. Addison-Wesley, Reading, Massachusetts, 1986.
- [3] Leslie Lamport. *L^AT_EX — A Document Preparation System*. Addison-Wesley, Reading, Massachusetts, second edition, 1994.

APPENDICES